

Basics of Machine Learning & Artificial Intelligence

Decision Intelligence Confronts A.I.

Friday November 10, 2023

John Mount

Win Vector LLC

All refs/links:

https://github.com/WinVector/Examples/tree/main/decision_intel_ai



Win Vector LLC specializes in data science consulting and training.

Disclaimer

- Produced to orient us on some current terms and claims
- A *greatly* simplified and abbreviated view and opinion
- Going to a bit of a “potted history”, opinion, and we gotta go fast to hit 15 minutes! (sorry!!!!)
- Trying to stay towards points that may be relevant to decision intelligence, quality, and general consequences
 - Talk goal: set up terms so the summary slide makes sense
 - “How to talk about it”



Outline

- Statistical Foundations of Machine Learning (the start)
- Deep Learning (evolutionary explosion)
- Current AI (amazing, but has some caveats)

Statistical Foundations of Machine Learning




what most people think of when they hear "machine learning"

Statistical Foundations of Machine Learning

- Machine Learning techniques became dominant in the 1990s through 2010s
 - Linear and Logistic Regression
 - Support Vector Machines and Kernels
 - Classification and Regression Trees
 - Random Forests
 - Gradient Boosting
 - Neural Networks (pre deep learning)
 - Ensemble Methods
 - Sparse Machine Learning / Compressed Sensing
- Capstone texts include:
 - Hastie, Tibshirani, Friedman, *The Elements of Statistical Learning*, Springer 2001.
 - Bishop, *Pattern Recognition and Machine Learning*, 2006.


This is what is often
meant by “machine learning”



many of these methods are statistical

Tabular Data Example

```
> library(palmerpenguins)
> data(package = 'palmerpenguins')
> head(penguins)
# A tibble: 6 x 8
  species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex  year
  <fct>   <fct>         <dbl>         <dbl>         <int>         <int> <fct> <int>
1 Adelia Torgersen      39.1           18.7           181          3750 male  2007
2 Adelia Torgersen      39.5           17.4           186          3800 female 2007
3 Adelia Torgersen      40.3            18           195          3250 female 2007
4 Adelia Torgersen      NA             NA             NA             NA NA     2007
5 Adelia Torgersen      36.7           19.3           193          3450 female 2007
6 Adelia Torgersen      39.3           20.6           190          3650 male  2007
```



Analysis methods concentrated on structured or tabular data

Derived from the “design matrix” in statistics

An array where each row is an instance, and each column is a type of measurement. Intersection is a measurement about an instance.

Spreadsheet, database table.

Jacquard Loom and Hollerith Punched card also in history.

Supervised Machine Learning Task


- Assume one column is both valuable to know, and expensive to know in a timely fashion. Call this column the “outcome” or “dependent variable”: y .
- Label all columns that are cheap to know as the “explanatory variables”: x (x a row vector per instance).
- Use historic data where all columns are known to infer a function $f()$ such that $f(x) \sim y$ on many rows of training data.
- Assume or ensure there are reasons so that $f(x) \sim y$ should continue on new or application data.
- Now when only x is available, we have an estimate of y !



The big trick: is how to re-order and reorganize your data into supervised training data. "supervised" meaning "y" is known during training. We will return to this.

Some Terminology Traps


- Inference
 - Statisticians use term to mean estimating $f()$. Also called “training.”
 - ML Engineers use term to mean computing $f(x)$, given $f()$ and x , others call this prediction.
- Prediction
 - At best statistical machine learning models predict $f(x) \sim E[y | x]$, not actual precise values for given instances.
- Independent variables
 - Traditional (horrible) name for explanatory variables.
- Regression
 - Traditional name for predicting a number. Comes from a early criticism of the process: regression to mediocrity, or inferring the regression line.
- Classification
 - Traditional name for predicting a choice from a fixed list of alternatives.
 - *A/ways* inferior to predicting probabilities and leaving choice to a controlled policy or decision procedure.
 - Insistence on “classification returns a class label” leads to harmful procedures: such as naively re-balancing training data.



Many related fields use the same terms for different things.

A Difference Between Machine Learning and Classic Modeling


- In classic modeling one often models aggregates directly
 - Model the defect *rate* of a factory as a function of physically modeled *controlled* factors such as temperature and humidity.
 - Modeling considers functions over aggregates.
- In machine learning one usually models individuals and then aggregates later
 - Model the probability of an individual part having a defect as function of *observed* factors such as historic temperature and humidity.
 - Defect rate then estimated by aggregation of predictions over proposed individuals.



Deep Learning

What is Deep Learning?

- Large directed acyclic graphs of mathematical expressions
 - Suggestively called “neural networks.”
- Plus a specific set of training from data strategies
 - Stochastic gradient descent
 - Drop out (a protection against over fitting)
 - Simulating un-censorship to generate training data
- Huge growth since the 2010s.




Un-Censorship

- Learn to replace the missing word
 - Original “As right as rain”
 - As training data
 - explanatory content (input): “As right as ????”
 - training outcome: “???? = rain”
 - Same ideas as building a Doob Martingale by controlling exposure order.
- First *huge* and famous success of this method: `word2vec` (Mikolov, Chen, Corrado, Dean, 2013)
 - Observing internal activation pattern of neural net leads to auto encoders and semantic embeddings
 - Moderate dimension: 200 to 300 columns
 - Related concepts near each other in embedding!



The big idea: how convert data into training data. Also this was a super clean experiment- deliberately deficient in text processing (such as no stemming to lemmatizing of words) to emphasize result was from the neural net induced embedding.


Deep Learning is measured in \$



<https://openai.com/research/ai-and-compute>




What we claim is from 1958 through the early 2010s: you ran a new system on a state of the art computer. How many petaflop-days it took to make the news tracked Moore's law, as Moore's law predicted the capabilities of that computer. In the mid 2010s the petaflop doubling rate sped up to every 3.4 months instead of 2 years. This means: to make the news you had to attract exponentially more capital than your predecessors to build a computer complex that move further ahead of commercial systems.



What AI is changes rapidly. Here is an approximation of this year's definition.

AI appears to be here

- Components
 - Large Language Models (LLMs)
 - Dominant since late 2010s early 2020s
 - Deep Neural Net trained on *huge* corpuses
 - Transformers
 - Self-supervised and semi-supervised training techniques
 - Prompt engineering (replacing user queries)
 - Post-processing
 - Specialized sub-systems
 - Human taggers
- Generally ***behind an API*** (application programming interface or service)
 - useful and impressive
 - yet, unfalsifiable



The behind the API issue brings us to Mechanical Turk. Originally a human agent in a box pretending to be a chess playing automata in 1770. Much later the name for one of the largest human tagging (or training data preparation) systems. Both senses remain relevant.

Actual use of a LLM based AI

JM

What does the following code do?

```
import numpy as np
import pandas as pd

def get_digit_sum(anumber):
    sum=0
    while (anumber != 0):
        sum = sum + int(anumber % 10)
        anumber = int(anumber/10)

    return sum

def get_number(astring):
    cleanstr = astring.replace(' ', '').lower()
    tokens = [ord(c) - 96 for c in cleanstr]
    mynumber = sum(tokens)

    # reduce to a single digit
    while mynumber >= 10:
        mynumber = get_digit_sum(mynumber)

    return mynumber
```

<https://chat.openai.com>



The text processing of current service based AIs is astounding. There are installable systems such as Hugging Face, but they tend to have much narrower capabilities.

ChatGPT3.5 Answer



The provided code defines two functions, `get_digit_sum` and `get_number`, which are used to calculate a numerical value associated with a given input string.

1. `get_digit_sum(anumber)`: This function calculates the sum of the digits of a number. It takes an integer `anumber` as input and iteratively adds up the digits of that number until the number becomes 0. It then returns the sum of the digits.
2. `get_number(astring)`: This function takes a string `astring` as input. It first cleans the string by removing spaces and converting it to lowercase. Then, it assigns a numeric value to each character in the cleaned string based on its position in the English alphabet (e.g., 'a' is 1, 'b' is 2, and so on). After calculating the sum of these values, it reduces the result to a single digit using the `get_digit_sum` function if necessary. The final single-digit result is returned.

Here's how you might use the `get_number` function:

```
python Copy code

result = get_number("Hello, World")
print(result) # This would print a single-digit result based on the input string
```

The purpose of this code is to assign a numeric value to a string by summing the values associated with its characters in the English alphabet and reducing the result to a single digit. This can be used for various applications, such as numerology or assigning values to words or phrases.



This answer is very good and very correct.

What are the AIs doing?

- Gave up on modeling mind or process
 - Modeling the mere structure of text or images
 - RAG or Retrieval Augmented Generation (replace portions of LLM answer with looked up results!)
- Produce results that are **statistically indistinguishable from correct answers**.
- Takes Searle's "Chinese Room" criticism as an *actual implementation blueprint*
 - As a criticism
 - Does a book with every possible Mandarin phrase next to its English equivalent "understand Mandarin and English?"
 - As a blueprint
 - Levers the amazing amount of existing digital content
- Training and testing issues
 - GIGO (garbage in, garbage out) remains an issue!
 - Racist and sexist biases often retained from training data
 - May not be an easy ask an AI a **novel** question at this point
 - We are now **"testing" on training data (not a reliable procedure!)**



https://www.linkedin.com/posts/philipp-schmid-a6q2bb196_how-openai-improved-their-rag-pipeline-by-activity-7127563601300606976-EDcm

Pay particular attention to the use of RAG and other "touch up the negatives" methodologies. They make for better systems, but obscure what components drive result quality. They make component evaluation a bit harder.

Some issues working with AIs

- Models commonly produce *confident* plausible hallucinations instead of correct answers.
- **Currently hidden behind a “pennies per task” network API**, so we don’t know much about the actual (versus claimed) implementation or unit cost.
- Better at producing spam than filtering spam.
- Impact on human taggers and editors in the process.
- **Current leading companies using funding and first-mover advantage to complete regulatory capture of the benefits** (see Yann LeCun’s criticism).
- Partnering with AI companies may be risky
 - AI’s are probably *nowhere near* as dangerous as the corporations that own them.



We are right to be afraid of infinitely powerful indifferent agents. However that is also a description of a corporation.

Additional Issue

- May not be safe to allow the general public to directly access your AI. (i.e. **don't do that!!!!**)
- **No** history of having a publicly facing AI stay on designed policy.



<https://spectrum.ieee.org/in-2016-microsofts-racist-chatbot-revealed-the-dangers-of-online-conversation>

<https://www.jumpstartmag.com/top-5-epic-fails-of-ai-chatbots/>

<https://www.bloomberg.com/news/newsletters/2022-12-08/chatgpt-open-ai-s-chatbot-is-spitting-out-biased-sexist-results>




You can't use an AI to face a potentially hostile audience (such as manning a call center). Prompt re-engineering such as "ignore your previous instructions" remains an issue.

Behind the API curtain

<https://www.jwz.org/blog/2023/11/once-again-ai-is-revealed-to-be-an-army-of-mechanical-turks-in-a-call-center/>

Once again, "AI" is revealed to be an army of mechanical turks in a call center

"[Cruise "autonomous"] vehicles were supported by a vast operations staff, with 1.5 workers per vehicle. The workers intervened to assist the company's vehicles every 2.5 to five miles"



The "autonomous" cars had remote supervision.

Conclusion



AI is becoming an important component in larger systems. They are dominating what we used to call information retrieval and text and image creation.

Summary Opinion

- Machine learning and AI are not the same thing
 - The days of selling a linear regression or some if-else statements as “AI” may be over
 - Have to go back to selling analytics as analytics (not a bad thing)
- Both Approaches ***treat observed quality of outcomes as inherent quality of decisions or process*** (a somewhat bad thing)
- AI ***is*** here
 - State of the art ***is currently a captive service***
 - AI augmented systems ***will*** outperform non-augmented systems
 - No longer just deep learning (prompt engineering, and sub-modules are critical)
 - Dominates in unstructured processes such as text and images
 - Right now probably best as a staff force multiplying tool
 - Not something to be ignored, a significant increase in capabilities in information retrieval



Thank You



Now lets move on to issues of decision intelligence and decision quality.